

pamica: GPU-accelerated Adaptive Mixture Independent Component Analysis in Python with Fortran parity

Seyed Yahya Shirazi¹, Arnaud Delorme^{1,2}, and Scott Makeig¹

¹ Swartz Center for Computational Neuroscience, Institute for Neural Computation, University of California San Diego, USA ² Centre de Recherche Cerveau et Cognition (CerCo), CNRS, University of Toulouse, France ¶ Corresponding author

DOI: 10.xxxxxx/draft

Software

- Review
- Repository
- Archive

Editor: Open Journals

Reviewers:

- @openjournals

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License (CC BY 4.0).

Summary

Independent Component Analysis (ICA) separates electroencephalographic and magnetoencephalographic (EEG/MEG) recordings into maximally independent sources that isolate brain, muscle, and artifact activities for downstream analysis (Iversen & Makeig, 2019; Makeig et al., 1995; Vigário et al., 1997). Adaptive Mixture ICA (AMICA) (Palmer et al., 2006, 2007, 2012) models each source with a flexible, self-adjusting probability density and lets several ICA models coexist, one per segment of a recording. Among the algorithms benchmarked by Delorme et al. (2012) it returned the least dependent and most dipolar decompositions, properties that study links to physiological interpretability. Its reference implementation, by Jason Palmer, is a Fortran program distributed as a compiled binary and usually run from MATLAB/EEGLAB; it runs only on the central processing unit (CPU) and has no Python interface.

pamica reproduces the reference Fortran results within numerical tolerance while running on the CPU, NVIDIA graphics processing units (GPUs, via CUDA), and Apple GPUs (through Apple's MLX array framework (Hannun et al., 2023)). It reimplements the algorithm on PyTorch (Paszke et al., 2019), NumPy (Harris et al., 2020), and SciPy (Virtanen et al., 2020), and exposes a scikit-learn-style estimator under a BSD-3-Clause license. It writes the format EEGLAB's AMICA loader reads, so established MATLAB tooling consumes its results unchanged. The software is at <https://github.com/scen/pAMICA> (archived at doi:10.5281/zenodo.21312148).

Statement of need

AMICA decompositions are well suited to equivalent-dipole source localization and automated component classification (Pion-Tonachini et al., 2019). Yet the reference implementation is used mainly through MATLAB, a growing obstacle as neuroimaging analysis moves toward Python, for example MNE-Python (Gramfort et al., 2013): modern pipelines need an AMICA that runs natively in Python and on a GPU, and that is validated to reproduce the Fortran reference numerically.

General-purpose Python ICA implementations do not fill this gap. scikit-learn and MNE-Python provide FastICA (Hyvärinen & Oja, 2000) and Infomax (Bell & Sejnowski, 1995; Lee et al., 1999), while Picard (Ablin et al., 2018) offers faster-converging maximum-likelihood ICA; none implement AMICA's mixture of models, adaptive generalized-Gaussian densities, or Newton updates, so none can reproduce its decompositions. pamica is for analysts who want AMICA-quality decompositions in Python, for anyone with a GPU who wants faster runs than the CPU-only binary, and for methodologists who need a transparent reference to build on.

41 The parity measurements reported here use EEG. The algorithm itself is modality-agnostic,
42 and an external user has fit rank-reduced, Maxwell-filtered MEG with it end to end, but parity
43 with the reference has not yet been measured on MEG.

44 State of the field

45 `pamica` is designed to work alongside the reference Fortran AMICA used with EEGLAB (Delorme
46 & Makeig, 2004): it keeps the same output format, adds a Python API with GPU support, and
47 can run the reference Fortran itself through a bundled dependency-free native build. Three
48 other Python AMICA reimplementations have appeared as of August 2026 (Esmaili, 2025;
49 Herforth, 2026; Huberty, 2025), oriented toward MNE-Python; `pamica` adds a scikit-learn-style
50 array API, output in EEGLAB's exact on-disk layout, an MLX backend, and an optional MNE-
51 Python wrapper. Its validation includes score functions checked against the literal Fortran
52 expressions to floating-point resolution, and a distributional comparison for the non-identifiable
53 multi-model case.

54 Software design

55 The central design decision was to define correctness as numerical parity with Palmer's Fortran,
56 and most of the other choices follow from it. Wrapping the binary would have secured parity but
57 inherited its CPU-only, MATLAB-facing design, while a reimplementations must demonstrate
58 parity. `pamica` ports the algorithm natively and also ships the reference binary as a runnable
59 engine, so users can check the parity claims on their own data and hardware. For the same
60 reason the port follows the reference's natural-gradient (Amari, 1998) expectation-maximization
61 (EM) formulation. An early Adam/autograd backend was removed because it converged to
62 different optima, which made results labeled "AMICA" hard to compare with the reference. The
63 port covers exact-EM mixture updates, a positive-definite Newton step (Palmer et al., 2008),
64 symmetric sphering, the five source-density families, a mixture of ICA models, component
65 sharing, and the mutual-information metrics used to score separation quality (Frank et al.,
66 2023).

67 Three array backends (PyTorch, MLX, NumPy) implement the same algorithm behind a
68 common estimator API, with the reference Fortran a fourth. NumPy is kept as a readable
69 executable specification. MLX is included because it is the more efficient path on Apple GPUs:
70 in the block-size sweep, PyTorch's Metal backend was slower than MLX at every size, and at
71 `pamica`'s shipped defaults slower than the CPU it runs beside (see the [documentation](#)). Double
72 precision is the default because the parity comparisons need it; single precision is available for
73 Apple GPUs, which have no float64.

74 Validation

75 Parity is measured two ways: by Hungarian-matched component correlation, and by the
76 Amari distance (Amari et al., 1996), a relabeling- and scale-invariant metric that needs no
77 assignment step. Both implementations ran AMICA's default 2000 iterations with Newton
78 disabled (`pamica`'s own default), to isolate the algorithm from initialization. With Newton
79 enabled and independent seeds, some of the weakest components settle into a different basin
80 of equal or higher likelihood: on one seed of three this affected ten of seventy components,
81 while the other two matched the reference at ~ 0.99 . A matched initialization restores ~ 0.997 ,
82 so this is a property of the initialization, not a parity defect. The single-model comparison uses
83 a well-determined external recording (NEMAR on002718, $k \approx 153$, where k = frames over
84 squared channel count (Frank et al., 2025)) alongside the bundled 32-channel sample ($k \approx 30$).
85 A mixture of ICA models is not partition-identifiable, so exact partition parity is the wrong
86 bar there; it is judged instead by whether the implementations sample a similar distribution of

87 solutions, over ensembles of 20 runs each (Figure 1). A permutation test finds no evidence
88 that cross-implementation agreement is worse than Fortran's own run-to-run agreement.

Table 1: Parity of pamica with the Fortran reference. Multi-model rows are over 20-run ensembles (190 within-, 400 cross-implementation pairs); sd is the standard deviation, given where computed. The final row is the one metric on which the ensembles differ significantly, and it is convergence speed rather than optimum quality: pamica reaches Fortran's mean by 200 iterations and passes it by 300.

Regime	Metric (dataset)	Result (mean)
Single	Log-likelihood gap (on002718)	within ~ 0.0005 of -3.6993
Single	Component correlation (on002718)	0.998
Single	Amari distance (bundled)	0.006
Single	Score functions, sufficient statistics	exact, $\sim 10^{-15}$
Multi	Correlation, one run: cross; within-Fortran	0.65; 0.64 (sd 0.05)
Multi	Amari, one run: cross; within-Fortran	0.163; 0.174 (sd 0.02)
Multi	Ensemble agreement, cross — within-Fortran	correlation +0.011 ($p = 0.96$); Amari -0.011 ($p > 0.999$)
Multi	Ensemble log-likelihood: Fortran; pamica	-3.3539 ; -3.3629 (Kolmogorov-Smirnov $p = 6 \times 10^{-5}$)

Multi-model AMICA (n_models=2): pamica vs Fortran ensembles, real sample EEG

Dashed vertical lines mark each distribution's mean.

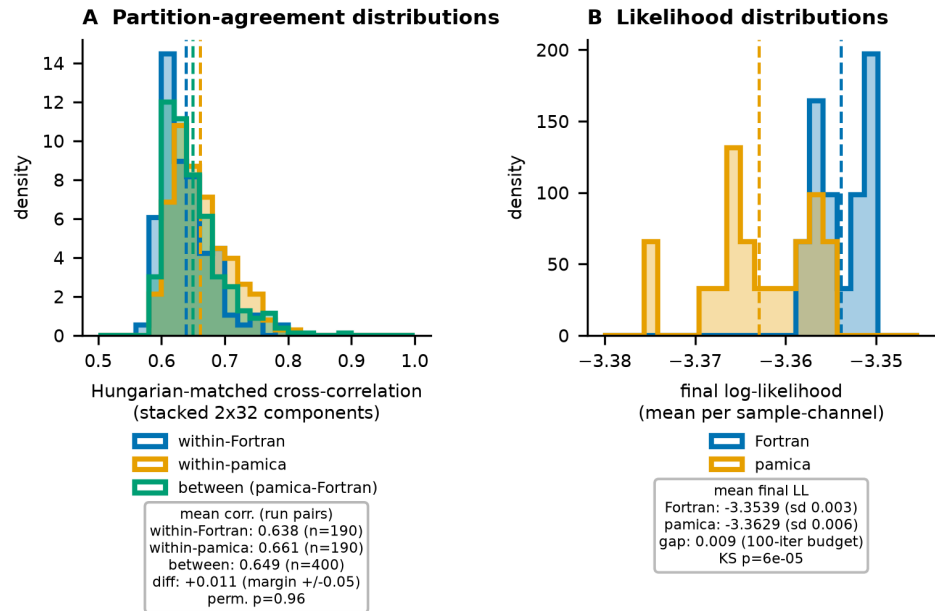


Figure 1: Multi-model ensemble partition-correlation (A) and log-likelihood (B) distributions, 20 pamica and 20 Fortran fits of the sample EEG; dashed lines mark each mean. A's three distributions overlap, so the single-run correlation reflects intrinsic run-to-run spread, not a gap to the reference. B's separation is Table 1's 0.009 gap on a ~ 0.035 axis.

89 All backends converge to the same single-model log-likelihood on real EEG (maximum pairwise
90 difference ~ 0.003), and single precision matches double to four or five significant digits on
91 that log-likelihood. Component-level float32 agreement is not yet characterized at a matched
92 iteration budget, so float64 remains the default for parity work, and double-precision CUDA is
93 the reproducible NVIDIA path.

94 On real 70-channel EEG at block_size=512, per-iteration cost is 25 ms for MLX on an Apple
95 GPU, 39 ms for double-precision CUDA on an RTX 4090, and 30 ms for native Fortran
96 on a 24-core i9-13900K, against 193 ms for PyTorch on an Apple-Silicon CPU and 255 ms
97 for PyTorch-Metal. That comparison is block-size-confounded: PyTorch-Metal is far more
98 block-size-sensitive than the CPU it runs beside, falling from 431 to 30.5 ms/iteration on
99 the bundled 32-channel sample between block_size=512 and pamica's current 8192 default,
100 still behind the CPU's 21.7 ms/iteration there, and to 13.5 ms/iteration, below the CPU's
101 15.8 ms, at a further-tuned single-block setting that is memory-limited rather than a free win,
102 since peak block memory scales with block_size, which is why 8192 stays the shipped default.
103 MLX stays the fastest Apple backend throughout the sweep. Full tables, the data-size sweep,
104 and reproduction commands are in the [documentation](#); the correctness harness never uses
105 synthetic data. The harness and sample data ship with the package, and every row of Table 1
106 but the two on002718 entries re-runs from the bundled sample alone.

Research impact statement

108 AMICA's source model and estimation were published in 2006-2008 (Palmer et al., 2006, 2007,
109 2008), and it ranked first of the twenty-two algorithms benchmarked by Delorme et al. (2012).
110 MNE-Python is publicly weighing which of the Python ports to adopt (mne-tools/mne-python

issue 13819), where an MNE core developer and the author of a competing port lean toward pamica, since this Center develops AMICA and maintains the reference Fortran and EEGLAB. sccn/pAMICA has been public since 2021, with public issues, pull requests, nine releases (six on PyPI), and a Zenodo archive. Three researchers outside the author group have filed issues, two mid-analysis on their own data: a 236-channel, 8.3-million-sample EEG decomposition (issue 207) and Maxwell-filtered 306-channel MEG (issue 221); the third is adopting AMICA for a multiverse analysis in MNE-Python (issue 206). Integration into this Center's Python preprocessing and into NEMAR (Delorme et al., 2022), the public archive we operate (800 datasets, ~40,000 participants, 55 TB), is in progress.

AI usage disclosure

Generative AI was used in this project, disclosed here under the journal's policy.

Tools. Anthropic's Claude models (Sonnet and Opus families), through the Claude Code command-line assistant. The instructions given to them are public in the repository (AGENTS.md, CLAUDE.md, .rules/).

Scope. The source code (translating the reference Fortran into Python, refactoring, scaffolding tests), the documentation, and the drafting and copy-editing of this manuscript.

Human oversight. The authors made the design decisions and take responsibility for the result. Parity as the correctness criterion, the natural-gradient EM formulation, the backend architecture, the distributional treatment of the multi-model case, and the acceptance thresholds were chosen by the authors. Every AI-assisted change was reviewed by a human before merge and run through the parity harness, which scores output against the reference binary on real recordings. The reported numbers came from running the software and were checked against their run records, as was every bibliographic entry.

Acknowledgments

We thank Jason Palmer and his advisor Ken Kreutz-Delgado, co-developers of AMICA, for the reference implementation, and the EEGLAB community for the tools and sample data used to validate this work. Two authors developed the methods pamica builds on: S.M. co-developed AMICA (Palmer et al., 2012) and A.D. is a lead developer of EEGLAB (Delorme & Makeig, 2004). This work was supported by The Swartz Foundation (Old Field, NY) to the Swartz Center for Computational Neuroscience and by National Institutes of Health grant R01-NS047293 (to A.D. and S.M.).

References

- Ablin, P., Cardoso, J.-F., & Gramfort, A. (2018). Faster independent component analysis by preconditioning with hessian approximations. *IEEE Transactions on Signal Processing*, 66(15), 4040–4053. <https://doi.org/10.1109/TSP.2018.2844203>
- Amari, S.-I. (1998). Natural gradient works efficiently in learning. *Neural Computation*, 10(2), 251–276. <https://doi.org/10.1162/089976698300017746>
- Amari, S.-I., Cichocki, A., & Yang, H. H. (1996). A new learning algorithm for blind signal separation. *Advances in Neural Information Processing Systems*, 8, 757–763.
- Bell, A. J., & Sejnowski, T. J. (1995). An information-maximization approach to blind separation and blind deconvolution. *Neural Computation*, 7(6), 1129–1159. <https://doi.org/10.1162/neco.1995.7.6.1129>

- 153 Delorme, A., & Makeig, S. (2004). EEGLAB: An open source toolbox for analysis of single-
154 trial EEG dynamics including independent component analysis. *Journal of Neuroscience*
155 *Methods*, 134(1), 9–21. <https://doi.org/10.1016/j.jneumeth.2003.10.009>
- 156 Delorme, A., Palmer, J., Onton, J., Oostenveld, R., & Makeig, S. (2012). Independent EEG
157 sources are dipolar. *PLoS ONE*, 7(2), e30135. [https://doi.org/10.1371/journal.pone.](https://doi.org/10.1371/journal.pone.0030135)
158 [0030135](https://doi.org/10.1371/journal.pone.0030135)
- 159 Delorme, A., Truong, D., Youn, C., Sivagnanam, S., Stirn, K., Yoshimoto, K., Poldrack,
160 R. A., Majumdar, A., & Makeig, S. (2022). NEMAR: An open access data, tools and
161 compute resource operating on neuroelectromagnetic data. *Database*, 2022, baac096.
162 <https://doi.org/10.1093/database/baac096>
- 163 Esmaeili, S. (2025). *Amica: Native python implementation of AMICA for MNE-Python and*
164 *scientific EEG workflows*. <https://github.com/snesmaeili/amica>.
- 165 Frank, G., Shirazi, S. Y., Palmer, J., Cauwenberghs, G., Makeig, S., & Delorme, A. (2023).
166 An exploration of optimal parameters for efficient blind source separation of EEG
167 recordings using AMICA. *2023 IEEE 23rd International Conference on Bioinformatics and*
168 *Bioengineering (BIBE)*, 205–210. <https://doi.org/10.1109/BIBE60311.2023.00040>
- 169 Frank, G., Shirazi, S. Y., Palmer, J., Cauwenberghs, G., Makeig, S., & Delorme, A. (2025).
170 Quantitative exploration of sufficient data quantities for EEG decomposition using AMICA.
171 *2025 12th International IEEE/EMBS Conference on Neural Engineering (NER)*, 532–536.
172 <https://doi.org/10.1109/NER61569.2025.11588993>
- 173 Gramfort, A., Luessi, M., Larson, E., Engemann, D. A., Strohmeier, D., Brodbeck, C., Goj,
174 R., Jas, M., Brooks, T., Parkkonen, L., & Hämäläinen, M. (2013). MEG and EEG data
175 analysis with MNE-Python. *Frontiers in Neuroscience*, 7, 267. <https://doi.org/10.3389/fnins.2013.00267>
- 176 Hannun, A., Digani, J., Katharopoulos, A., & Collobert, R. (2023). *MLX: Efficient and flexible*
177 *machine learning on apple silicon*. <https://github.com/ml-explore/mlx>.
- 179 Harris, C. R., Millman, K. J., Walt, S. J. van der, Gommers, R., Virtanen, P., Cournapeau, D.,
180 Wieser, E., Taylor, J., Berg, S., Smith, N. J., & others. (2020). Array programming with
181 NumPy. *Nature*, 585(7825), 357–362. <https://doi.org/10.1038/s41586-020-2649-2>
- 182 Herforth, J. (2026). *Pyamica: PyTorch AMICA, adaptive mixture independent component*
183 *analysis*. <https://github.com/DerAndereJohannes/pyamica>.
- 184 Huberty, S. (2025). *Amica-python: Python implementation of adaptive mixture ICA*. <https://github.com/scott-huberty/amica-python>.
- 185 Hyvärinen, A., & Oja, E. (2000). Independent component analysis: Algorithms and applications.
186 *Neural Networks*, 13(4-5), 411–430. [https://doi.org/10.1016/S0893-6080\(00\)00026-5](https://doi.org/10.1016/S0893-6080(00)00026-5)
- 187 Iversen, J. R., & Makeig, S. (2019). MEG/EEG data analysis using EEGLAB. In
188 *Magnetoencephalography: From signals to dynamic cortical networks* (pp. 391–406).
189 Springer International Publishing. https://doi.org/10.1007/978-3-319-62657-4_8-1
- 190 Lee, T.-W., Girolami, M., & Sejnowski, T. J. (1999). Independent component analysis using
191 an extended infomax algorithm for mixed subgaussian and supergaussian sources. *Neural*
192 *Computation*, 11(2), 417–441. <https://doi.org/10.1162/089976699300016719>
- 193 Makeig, S., Bell, A. J., Jung, T.-P., & Sejnowski, T. J. (1995). Independent component
194 analysis of electroencephalographic data. *Advances in Neural Information Processing*
195 *Systems*, 8, 145–151.
- 196 Palmer, J. A., Kreutz-Delgado, K., & Makeig, S. (2006). Super-Gaussian mixture source model
197 for ICA. *Independent Component Analysis and Blind Signal Separation: Proceedings of the*
198 *6th International Conference (ICA 2006)*, 854–861. https://doi.org/10.1007/11679363_106

- 200 Palmer, J. A., Kreutz-Delgado, K., & Makeig, S. (2012). *AMICA: An adaptive*
201 *mixture of independent component analyzers with shared components*. Swartz
202 Center for Computational Neuroscience, University of California San Diego.
203 https://sccn.ucsd.edu/~jason/amica_a.pdf
- 204 Palmer, J. A., Kreutz-Delgado, K., Rao, B. D., & Makeig, S. (2007). Modeling and estimation
205 of dependent subspaces with non-radially symmetric and skewed densities. In M. E. Davies,
206 C. J. James, S. A. Abdallah, & M. D. Plumbley (Eds.), *Independent component analysis*
207 *and signal separation: Proceedings of the 7th international conference (ICA 2007)* (pp.
208 97–104). Springer Berlin Heidelberg. https://doi.org/10.1007/978-3-540-74494-8_13
- 209 Palmer, J. A., Kreutz-Delgado, K., Rao, B. D., & Makeig, S. (2008). Newton method for the
210 ICA mixture model. *2008 IEEE International Conference on Acoustics, Speech and Signal*
211 *Processing (ICASSP)*, 1805–1808. <https://doi.org/10.1109/ICASSP.2008.4517982>
- 212 Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin,
213 Z., Gimelshein, N., Antiga, L., & others. (2019). PyTorch: An imperative style, high-
214 performance deep learning library. *Advances in Neural Information Processing Systems*, 32,
215 8024–8035.
- 216 Pion-Tonachini, L., Kreutz-Delgado, K., & Makeig, S. (2019). ICLabel: An automated
217 electroencephalographic independent component classifier, dataset, and website.
218 *NeuroImage*, 198, 181–197. <https://doi.org/10.1016/j.neuroimage.2019.05.026>
- 219 Vigário, R., Jousmäki, V., Härmäläinen, M., Hari, R., & Oja, E. (1997). Independent component
220 analysis for identification of artifacts in magnetoencephalographic recordings. *Advances in*
221 *Neural Information Processing Systems*, 10, 229–235.
- 222 Virtanen, P., Gommers, R., Oliphant, T. E., Haberland, M., Reddy, T., Cournapeau, D.,
223 Burovski, E., Peterson, P., Weckesser, W., Bright, J., & others. (2020). SciPy 1.0:
224 Fundamental algorithms for scientific computing in python. *Nature Methods*, 17(3),
225 261–272. <https://doi.org/10.1038/s41592-019-0686-2>