

pamica: GPU-accelerated Adaptive Mixture Independent Component Analysis in Python with Fortran parity

Seyed Yahya Shirazi¹, Arnaud Delorme^{1,2}, and Scott Makeig¹

¹ Swartz Center for Computational Neuroscience, Institute for Neural Computation, University of California San Diego, USA ² Centre de Recherche Cerveau et Cognition (CerCo), CNRS, University of Toulouse, France ¶ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [Open Journals](#)

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Independent Component Analysis (ICA) separates electroencephalographic and magnetoencephalographic (EEG/MEG) recordings into maximally independent sources that isolate brain, muscle, and artifact activities for downstream analysis (Iversen & Makeig, 2019; Makeig et al., 1995; Vigário et al., 1997). Adaptive Mixture ICA (AMICA) (Palmer et al., 2006, 2007, 2012) models each source with a flexible, self-adjusting probability density and lets several ICA models coexist, one per segment of a recording. Among the algorithms benchmarked by Delorme et al. (2012) it returned the least dependent and most dipolar decompositions, properties that the study links to physiological interpretability. Its reference implementation, by Jason Palmer, is a Fortran program distributed as a compiled binary and usually run from MATLAB/EEGLAB; it runs only on the central processing unit (CPU) and has no Python interface.

pamica reproduces the reference Fortran's single-model results within numerical tolerance, and its multi-model solutions spread like the reference's own, while running on the CPU, NVIDIA graphics processing units (GPUs, via CUDA), and Apple GPUs (through Apple's MLX array framework (Hannun et al., 2023)). It reimplements the algorithm on PyTorch (Paszke et al., 2019), NumPy (Harris et al., 2020), and SciPy (Virtanen et al., 2020), and exposes a scikit-learn-style estimator under a BSD-3-Clause license. It writes the format EEGLAB's AMICA loader reads, so established MATLAB tooling consumes its results unchanged. The software is at <https://github.com/sccn/pAMICA> (archived at doi:10.5281/zenodo.21312148).

Statement of need

AMICA decompositions are well suited to equivalent-dipole source localization and automated component classification (Pion-Tonachini et al., 2019). Yet the reference implementation is used mainly through MATLAB, a growing obstacle as neuroimaging analysis moves toward Python, for example MNE-Python (Gramfort et al., 2013): modern pipelines need an AMICA that runs natively in Python and on a GPU, and that is validated to reproduce the Fortran reference numerically.

General-purpose Python ICA implementations do not fill this gap. scikit-learn and MNE-Python provide FastICA (Hyvärinen & Oja, 2000) and Infomax (Bell & Sejnowski, 1995; Lee et al., 1999), while Picard (Ablin et al., 2018) offers faster-converging maximum-likelihood ICA; none implement AMICA's mixture of models, adaptive generalized-Gaussian densities, or Newton updates, so none can reproduce its decompositions. pamica serves analysts who want AMICA-quality decompositions in Python and anyone with a GPU who wants faster runs than

the CPU-only binary; it is also a transparent reference for methodologists to build on. The parity measurements reported here use EEG. The algorithm itself is modality-agnostic, and an external user has fit rank-reduced, Maxwell-filtered MEG with it end to end (see the research impact statement), but parity with the reference has not yet been measured on MEG.

State of the field

pamica is designed to work alongside the reference Fortran AMICA used with EEGLAB (Delorme & Makeig, 2004): it keeps the same output format, adds a Python API with GPU support, and can run the reference Fortran itself through a bundled dependency-free native build. Three other Python AMICA reimplementations have appeared as of August 2026 (Esmaili, 2025; Herforth, 2026; Huberty, 2025), oriented toward MNE-Python; pamica adds a scikit-learn-style array API, output in EEGLAB's exact on-disk layout, an MLX backend, and an optional MNE-Python wrapper. Its validation includes score functions checked against the literal Fortran expressions to floating-point resolution, and a distributional comparison for the non-identifiable multi-model case.

Software design

The central design decision was to define correctness as numerical parity with Palmer's Fortran, and most of the other choices follow from it. Wrapping the binary would have secured parity but inherited its CPU-only, MATLAB-facing design, while a reimplementation must demonstrate parity. pamica ports the algorithm natively and also ships the reference binary as a runnable engine, so users can check the parity claims on their own data and hardware. For the same reason the port follows the reference's natural-gradient (Amari, 1998) expectation-maximization (EM) formulation. An early Adam/autograd backend was removed because it converged to different optima, which made results labeled "AMICA" hard to compare with the reference. The port covers exact-EM mixture updates, a positive-definite Newton step (Palmer et al., 2008), symmetric sphering, the five source-density families, a mixture of ICA models, component sharing, and the mutual-information metrics used to score separation quality (Frank et al., 2023).

Three array backends (PyTorch, MLX, NumPy) implement the same algorithm behind a common estimator API, with the reference Fortran a fourth. NumPy is kept as a readable executable specification. MLX is included because it is the more efficient path on Apple GPUs: in the block-size sweep, PyTorch's Metal backend was slower than MLX at every size, and at pamica's shipped defaults slower than the CPU it runs beside (see the documentation). Double precision is the default because the parity comparisons need it; single precision is available for Apple GPUs, which have no float64.

Validation

Parity is measured two ways: by Hungarian-matched component correlation, and by the Amari distance (Amari et al., 1996), a relabeling- and scale-invariant metric that needs no assignment step. Both implementations ran EEGLAB's default of 2000 iterations with Newton disabled (pamica's own default), to isolate the algorithm from initialization. With Newton enabled, some of the weakest components settle into different basins of equal or higher likelihood from different starts, in the reference's runs as in pamica's: against one reference run, one pamica seed of three differed on eight of seventy components and the other two matched at 0.995 and 0.996, and two reference runs from different seeds differed on two. From a shared start the two implementations end with the same components (correlation 0.9999998). The single-model comparison uses a well-determined external recording (OpenNeuro ds002718, mirrored in the NEMAR archive (Delorme et al., 2022) as on002718, $k \approx 153$, where k = frames over squared

channel count (Frank et al., 2025)) alongside the bundled 32-channel sample ($k \approx 30$). On the bundled sample one of the five reference runs ended in a lower-likelihood basin and sets the five-pair Amari mean; over 50 run pairs the distance between the implementations (0.004) is close to the reference's own run-to-run distance (0.005). A mixture of ICA models is not partition-identifiable, so the implementations are compared by the distributions of solutions they sample, over ensembles of 20 runs each (Figure 1). One-sided run-level permutation tests of whether cross-implementation agreement is worse than the reference's own run-to-run agreement give $p = 0.88$ by correlation and $p = 0.051$ by Amari distance, on which pamica's ensemble spreads slightly more than the reference's.

Table 1: Parity of pamica with the Fortran reference. Multi-model rows are over 20-run ensembles (190 within-Fortran and 190 within-pamica, plus 400 cross-implementation pairs); sd is the standard deviation, given where computed; p values are one-sided run-level permutation tests.

Regime	Metric (dataset)	Result (mean)
Single	Log-likelihood gap (on002718)	6×10^{-6} , both at -3.6993
Single	Component correlation (on002718)	0.9996 (Fortran vs Fortran 0.998)
Single	Amari distance (bundled)	0.011 (5 run pairs); 0.004 (50 pairs)
Single	Score functions; sufficient statistics	$\leq 2 \times 10^{-15}$ absolute; 3×10^{-16} relative
Multi	Correlation, one run: cross; within-Fortran	0.632; 0.626 (sd 0.04)
Multi	Amari, one run: cross; within-Fortran	0.172; 0.166 (sd 0.02)
Multi	Ensemble agreement, cross — within-Fortran	correlation +0.006 ($p = 0.88$); Amari +0.005 ($p = 0.051$)
Multi	Ensemble log-likelihood: Fortran; pamica	-3.3543 ; -3.3541 (Kolmogorov-Smirnov $p = 0.83$)

Multi-model AMICA (n_models=2): pamica vs Fortran ensembles, real sample EEG

Vertical lines mark each distribution's mean, in its line style.




Figure 1: Multi-model ensemble partition-correlation (A) and log-likelihood (B) distributions, 20 pamica and 20 Fortran fits of the sample EEG; vertical lines mark each mean. A's three distributions overlap: the single-run correlation of ~ 0.63 matches the reference's agreement with itself. B's two means lie 0.0002 apart on a ~ 0.02 axis.

From a shared start with the harness settings, the single-precision MLX backend and the double-precision PyTorch backend differ by 5×10^{-6} in log-likelihood after 100 iterations, with mean matched component correlation 0.9999999 and Amari distance 5×10^{-5} . In the throughput benchmark's runs on a 30,000-frame excerpt of the external recording, the backends that share a start agree after 25 iterations to 10^{-5} in log-likelihood on 32 and 48 channels; with 70 channels ($k \approx 6$) they differ by up to 10^{-3} , about the effect of perturbing one sample by 10^{-9} μ V. Double precision remains the default for parity work, and double-precision CUDA is the reproducible NVIDIA path.

On real 70-channel EEG at block_size=512, per-iteration cost is 25 ms for MLX on an Apple GPU, 39 ms for double-precision CUDA on an RTX 4090, and 30 ms for native Fortran on a 24-core i9-13900K, against 193 ms for PyTorch on an Apple-Silicon CPU and 255 ms for PyTorch-Metal. That comparison is block-size-confounded: PyTorch-Metal is far more block-size-sensitive than the CPU it runs beside, falling from 431 to 30.5 ms/iteration on the bundled 32-channel sample between block_size=512 and pamica's current 8192 default, still behind the CPU's 21.7 ms/iteration there, and to 13.5 ms/iteration, below the CPU's 15.8 ms, at a further-tuned single-block setting. Peak block memory grows with block_size, which is why 8192 stays the shipped default. MLX stays the fastest Apple backend throughout the sweep. Full tables, the data-size sweep, and reproduction commands are in the [documentation](#); the correctness harness never uses synthetic data. The harness and sample data are in the source repository, and every row of Table 1 but the two on002718 entries re-runs from the bundled sample alone.

117 Research impact statement

118 AMICA's source model and estimation were published in 2006-2008 (Palmer et al., 2006, 2007,
119 2008), and it ranked first of the twenty-two algorithms benchmarked by Delorme et al. (2012).
120 MNE-Python is publicly weighing which of the Python ports to adopt (mne-tools/mne-python
121 issue 13819), where an MNE core developer and the author of a competing port lean toward
122 pamica, since this Center develops AMICA and maintains the reference Fortran and EEGLAB.

123 sccn/pAMICA has been public since 2021, with public issues, pull requests, nine releases (six
124 on PyPI), and a Zenodo archive. Three researchers outside the author group have filed
125 issues, two of them mid-analysis on their own data: a 236-channel, 8.3-million-sample EEG
126 decomposition (issue 207) and Maxwell-filtered 306-channel MEG (issue 221); the third is
127 adopting AMICA for a multiverse analysis in MNE-Python (issue 206). Integration into this
128 Center's Python preprocessing and into NEMAR, the public archive we operate (800 datasets,
129 ~40,000 participants, 55 TB), is in progress.

130 AI usage disclosure

131 Generative AI was used in this project, disclosed here under the journal's policy.

132 **Tools.** Anthropic's Claude models (Sonnet and Opus families), through the Claude Code
133 command-line assistant, with the Research Skills plugin collection ([https://github.com/
134 neuromechanist/research-skills](https://github.com/neuromechanist/research-skills)) (Shirazi, 2026) as the working harness, which Claude Code
135 does not provide on its own. Its project plugin supplies the project's design and planning
136 workflow, the repository instructions (AGENTS.md, CLAUDE.md, .rules/), the phased, issue-
137 tracked development and the pull-request review; its manuscript plugin supplies the manuscript
138 review. The instructions given to the models are public in the repository.

139 **Scope.** The source code (translating the reference Fortran into Python, refactoring, scaffolding
140 tests), the documentation, and the drafting and copy-editing of this manuscript.

141 **Human oversight.** The authors made the design decisions and take responsibility for the
142 result. Parity as the correctness criterion, the natural-gradient EM formulation, the backend
143 architecture, the distributional treatment of the multi-model case, and the acceptance thresholds
144 were chosen by the authors. Every AI-assisted change was reviewed by a human before merge
145 and run through the parity harness, which scores output against the reference binary on real
146 recordings. The reported numbers came from running the software and were checked against
147 their run records, as was every bibliographic entry.

148 Acknowledgments

149 We thank Jason Palmer and his advisor Ken Kreutz-Delgado, co-developers of AMICA, for
150 the reference implementation, and the EEGLAB community for the tools and sample data
151 used to validate this work. Two authors developed the methods pamica builds on: S.M.
152 co-developed AMICA (Palmer et al., 2012) and A.D. is a lead developer of EEGLAB (Delorme
153 & Makeig, 2004). This work was supported by The Swartz Foundation (Old Field, NY) to
154 the Swartz Center for Computational Neuroscience and by National Institutes of Health grant
155 R01-NS047293 (to A.D. and S.M.).

156 References

157 Ablin, P., Cardoso, J.-F., & Gramfort, A. (2018). Faster independent component analysis
158 by preconditioning with hessian approximations. *IEEE Transactions on Signal Processing*,
159 66(15), 4040–4053. <https://doi.org/10.1109/TSP.2018.2844203>

- Amari, S.-I. (1998). Natural gradient works efficiently in learning. *Neural Computation*, 10(2), 251–276. <https://doi.org/10.1162/089976698300017746>
- Amari, S.-I., Cichocki, A., & Yang, H. H. (1996). A new learning algorithm for blind signal separation. *Advances in Neural Information Processing Systems*, 8, 757–763.
- Bell, A. J., & Sejnowski, T. J. (1995). An information-maximization approach to blind separation and blind deconvolution. *Neural Computation*, 7(6), 1129–1159. <https://doi.org/10.1162/neco.1995.7.6.1129>
- Delorme, A., & Makeig, S. (2004). EEGLAB: An open source toolbox for analysis of single-trial EEG dynamics including independent component analysis. *Journal of Neuroscience Methods*, 134(1), 9–21. <https://doi.org/10.1016/j.jneumeth.2003.10.009>
- Delorme, A., Palmer, J., Onton, J., Oostenveld, R., & Makeig, S. (2012). Independent EEG sources are dipolar. *PLoS ONE*, 7(2), e30135. <https://doi.org/10.1371/journal.pone.0030135>
- Delorme, A., Truong, D., Youn, C., Sivagnanam, S., Stirn, K., Yoshimoto, K., Poldrack, R. A., Majumdar, A., & Makeig, S. (2022). NEMAR: An open access data, tools and compute resource operating on neuroelectromagnetic data. *Database*, 2022, baac096. <https://doi.org/10.1093/database/baac096>
- Esmaeili, S. (2025). *Amica: Native python implementation of AMICA for MNE-Python and scientific EEG workflows*. <https://github.com/snesmaeili/amica>.
- Frank, G., Shirazi, S. Y., Palmer, J., Cauwenberghs, G., Makeig, S., & Delorme, A. (2023). An exploration of optimal parameters for efficient blind source separation of EEG recordings using AMICA. *2023 IEEE 23rd International Conference on Bioinformatics and Bioengineering (BIBE)*, 205–210. <https://doi.org/10.1109/BIBE60311.2023.00040>
- Frank, G., Shirazi, S. Y., Palmer, J., Cauwenberghs, G., Makeig, S., & Delorme, A. (2025). Quantitative exploration of sufficient data quantities for EEG decomposition using AMICA. *2025 12th International IEEE/EMBS Conference on Neural Engineering (NER)*, 532–536. <https://doi.org/10.1109/NER61569.2025.11588993>
- Gramfort, A., Luessi, M., Larson, E., Engemann, D. A., Strohmeier, D., Brodbeck, C., Goj, R., Jas, M., Brooks, T., Parkkonen, L., & Hämäläinen, M. (2013). MEG and EEG data analysis with MNE-Python. *Frontiers in Neuroscience*, 7, 267. <https://doi.org/10.3389/fnins.2013.00267>
- Hannun, A., Digani, J., Katharopoulos, A., & Collobert, R. (2023). *MLX: Efficient and flexible machine learning on apple silicon*. <https://github.com/ml-explore/mlx>.
- Harris, C. R., Millman, K. J., Walt, S. J. van der, Gommers, R., Virtanen, P., Cournapeau, D., Wieser, E., Taylor, J., Berg, S., Smith, N. J., & others. (2020). Array programming with NumPy. *Nature*, 585(7825), 357–362. <https://doi.org/10.1038/s41586-020-2649-2>
- Herforth, J. (2026). *Pyamica: PyTorch AMICA, adaptive mixture independent component analysis*. <https://github.com/DerAndereJohannes/pyamica>.
- Huberty, S. (2025). *Amica-python: Python implementation of adaptive mixture ICA*. <https://github.com/scott-huberty/amica-python>.
- Hyvärinen, A., & Oja, E. (2000). Independent component analysis: Algorithms and applications. *Neural Networks*, 13(4-5), 411–430. [https://doi.org/10.1016/S0893-6080\(00\)00026-5](https://doi.org/10.1016/S0893-6080(00)00026-5)
- Iversen, J. R., & Makeig, S. (2019). MEG/EEG data analysis using EEGLAB. In *Magnetoencephalography: From signals to dynamic cortical networks* (pp. 391–406). Springer International Publishing. https://doi.org/10.1007/978-3-319-62657-4_8-1
- Lee, T.-W., Girolami, M., & Sejnowski, T. J. (1999). Independent component analysis using an extended infomax algorithm for mixed subgaussian and supergaussian sources. *Neural*

- 207 *Computation*, 11(2), 417–441. <https://doi.org/10.1162/089976699300016719>
- 208 Makeig, S., Bell, A. J., Jung, T.-P., & Sejnowski, T. J. (1995). Independent component
209 analysis of electroencephalographic data. *Advances in Neural Information Processing*
210 *Systems*, 8, 145–151.
- 211 Palmer, J. A., Kreutz-Delgado, K., & Makeig, S. (2006). Super-Gaussian mixture source model
212 for ICA. *Independent Component Analysis and Blind Signal Separation: Proceedings of the*
213 *6th International Conference (ICA 2006)*, 854–861. https://doi.org/10.1007/11679363_106
- 214 Palmer, J. A., Kreutz-Delgado, K., & Makeig, S. (2012). *AMICA: An adaptive*
215 *mixture of independent component analyzers with shared components*. Swartz
216 Center for Computational Neuroscience, University of California San Diego.
217 https://sccn.ucsd.edu/~jason/amica_a.pdf
- 218 Palmer, J. A., Kreutz-Delgado, K., Rao, B. D., & Makeig, S. (2007). Modeling and estimation
219 of dependent subspaces with non-radially symmetric and skewed densities. In M. E. Davies,
220 C. J. James, S. A. Abdallah, & M. D. Plumbley (Eds.), *Independent component analysis*
221 *and signal separation: Proceedings of the 7th international conference (ICA 2007)* (pp.
222 97–104). Springer Berlin Heidelberg. https://doi.org/10.1007/978-3-540-74494-8_13
- 223 Palmer, J. A., Kreutz-Delgado, K., Rao, B. D., & Makeig, S. (2008). Newton method for the
224 ICA mixture model. *2008 IEEE International Conference on Acoustics, Speech and Signal*
225 *Processing (ICASSP)*, 1805–1808. <https://doi.org/10.1109/ICASSP.2008.4517982>
- 226 Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin,
227 Z., Gimelshein, N., Antiga, L., & others. (2019). PyTorch: An imperative style, high-
228 performance deep learning library. *Advances in Neural Information Processing Systems*, 32,
229 8024–8035.
- 230 Pion-Tonachini, L., Kreutz-Delgado, K., & Makeig, S. (2019). ICLabel: An automated
231 electroencephalographic independent component classifier, dataset, and website.
232 *NeuroImage*, 198, 181–197. <https://doi.org/10.1016/j.neuroimage.2019.05.026>
- 233 Shirazi, S. Y. (2026). *Research skills*. <https://github.com/neuromechanist/research-skills>.
234 <https://doi.org/10.5281/zenodo.20696515>
- 235 Vigário, R., Jousmäki, V., Hämäläinen, M., Hari, R., & Oja, E. (1997). Independent component
236 analysis for identification of artifacts in magnetoencephalographic recordings. *Advances in*
237 *Neural Information Processing Systems*, 10, 229–235.
- 238 Virtanen, P., Gommers, R., Oliphant, T. E., Haberland, M., Reddy, T., Cournapeau, D.,
239 Burovski, E., Peterson, P., Weckesser, W., Bright, J., & others. (2020). SciPy 1.0:
240 Fundamental algorithms for scientific computing in python. *Nature Methods*, 17(3),
241 261–272. <https://doi.org/10.1038/s41592-019-0686-2>